

ICML 2001

Conditional Random Fields: Probabilistic Models for Segmenting and Labeling Sequence Data

John Lafferty, Andrew McCallum, Fernando Pereira

Presentation by Rongkun Shen

Nov. 20, 2003

Sequence Segmenting and Labeling

- Goal: mark up sequences with content tags
- Application in computational biology
 - DNA and protein sequence alignment
 - Sequence homolog searching in databases
 - Protein secondary structure prediction
 - RNA secondary structure analysis
- Application in computational linguistics & computer science
 - Text and speech processing, including topic segmentation, part-of-speech (POS) tagging
 - Information extraction
 - Syntactic disambiguation

Example: Protein secondary structure prediction

Conf: 977621015677468999723631357600330223342057899861488356412238

Pred: CCCCCCCCCCCCEEEEECCCCCCCCCCCCHHHHHHHHHHHHCCCEEEHHCC

AA: EKKSINECDLKGKKVLIRVDFNVPVKNGKITNDYRIRSALPTLKKVLTEGGSCVLMSHLG

10 20 30 40 50 60

Conf: 85576422245412347898510001047899999874033445740023666631258

Pred: CCCCCCCCCCCCCCCCCCCCCCCCCCCCHHHHHHHHHHHCCCCCCCCCCCHHHHHHCC

AA: RPKGIPMAQAGKIRSTGGVPGFQQKATLKPVAKRLSELLLRPVTFAPDCLNAADVVSKMS

70 80 90 100 110 120

Conf: 874688611002343044310017899999875053355212244334552001322452

Pred: CCCEEEECCCHHHHHCCCCCHHHHHHHHHCCCEEEECCCCCCCCCCCCCHHHH


AA: PGDVVLLENVRFYKEEGSKKAKDREAMAKILASYGDVYISDAFGTAHRDSATMTGIPKIL

130 140 150 160 170 180

Generative Models

- Hidden Markov models (HMMs) and stochastic grammars
 - Assign a joint probability to paired observation and label sequences
 - The parameters typically trained to maximize the joint likelihood of train examples

Standard tool is the hidden Markov Model (HMM).



$$P(\mathbf{X}, \mathbf{Y}) = \prod_i P(\mathbf{X}_i | \mathbf{Y}_i) P(\mathbf{Y}_i | \mathbf{Y}_{i-1})$$

Generative Models (cont'd)

- Difficulties and disadvantages
 - Need to enumerate all possible observation sequences
 - Not practical to represent multiple interacting features or long-range dependencies of the observations
 - Very strict independence assumptions on the observations


Conditional Models

- Conditional probability $P(\text{label sequence } \mathbf{y} / \text{observation sequence } \mathbf{x})$ rather than joint probability $P(\mathbf{y}, \mathbf{x})$
 - Specify the probability of possible label sequences given an observation sequence
- Allow arbitrary, non-independent features on the observation sequence \mathbf{X}
- The probability of a transition between labels may depend on **past** and **future** observations
 - Relax strong independence assumptions in generative models

Discriminative Models

Maximum Entropy Markov Models (MEMMs)

- Exponential model
- Given training set X with label sequence Y :
 - Train a model θ that maximizes $P(Y|X, \theta)$
 - For a new data sequence x , the predicted label y maximizes $P(y|x, \theta)$
 - Notice the per-state normalization




$$P(y' | y, x) = \frac{1}{Z(y, x)} \exp \left(\sum_k \underbrace{\lambda_k}_{\text{weight}} \underbrace{f_k(x, y, y')}_{\text{feature}} \right)$$

MEMMs (cont'd)

- MEMMs have all the advantages of Conditional Models
- Per-state normalization: all the mass that arrives at a state must be distributed among the possible successor states (“conservation of score mass”)
- Subject to Label Bias Problem
 - Bias toward states with fewer outgoing transitions

Label Bias Problem


- Consider this MEMM:



- $P(1 \text{ and } 2 \mid ro) = P(2 \mid 1 \text{ and } ro)P(1 \mid ro) = P(2 \mid 1 \text{ and } o)P(1 \mid r)$
 $P(1 \text{ and } 2 \mid ri) = P(2 \mid 1 \text{ and } ri)P(1 \mid ri) = P(2 \mid 1 \text{ and } i)P(1 \mid r)$
- Since $P(2 \mid 1 \text{ and } x) = 1$ for all x , $P(1 \text{ and } 2 \mid ro) = P(1 \text{ and } 2 \mid ri)$
In the training data, label value 2 is the only label value observed after label value 1
Therefore $P(2 \mid 1) = 1$, so $P(2 \mid 1 \text{ and } x) = 1$ for all x
- However, we expect $P(1 \text{ and } 2 \mid ri)$ to be greater than $P(1 \text{ and } 2 \mid ro)$.
- Per-state normalization does not allow the required expectation

Solve the Label Bias Problem

- Change the state-transition structure of the model




- Not always practical to change the set of states
- Start with a fully-connected model and let the training procedure figure out a good structure
 - Preclude the use of prior, which is very valuable (e.g. in information extraction)

Random Field

Let $G = (Y, E)$ be a graph where each vertex Y_v is a random variable
Suppose $P(Y_v | \text{all other } Y) = P(Y_v | \text{neighbors}(Y_v))$ then Y is a
random field

Example:



- $P(Y_5 | \text{all other } Y) = P(Y_5 | Y_4, Y_6)$

Conditional Random Fields (CRFs)

- CRFs have all the advantages of MEMMs without label bias problem
 - MEMM uses **per-state exponential** model for the conditional probabilities of next states given the **current state**
 - CRF has a **single exponential** model for the joint probability of the entire sequence of labels given the **observation sequence**
- Undirected acyclic graph
- Allow some transitions “vote” more strongly than others depending on the corresponding observations

Definition of CRFs


\mathbf{X} is a random variable over data sequences to be labeled

\mathbf{Y} is a random variable over corresponding label sequences

Definition. Let $G = (V, E)$ be a graph such that $\mathbf{Y} = (\mathbf{Y}_v)_{v \in V}$, so that \mathbf{Y} is indexed by the vertices of G . Then (\mathbf{X}, \mathbf{Y}) is a *conditional random field* in case, when conditioned on \mathbf{X} , the random variables \mathbf{Y}_v obey the Markov property with respect to the graph: $p(\mathbf{Y}_v | \mathbf{X}, \mathbf{Y}_w, w \neq v) = p(\mathbf{Y}_v | \mathbf{X}, \mathbf{Y}_w, w \sim v)$, where $w \sim v$ means that w and v are neighbors in G .

Example of CRFs

Suppose $P(Y_v | X, \text{all other } Y) = P(Y_v | X, \text{neighbors}(Y_v))$
then X with Y is a **conditional random field**



- $P(Y_3 | X, \text{all other } Y) = P(Y_3 | X, Y_2, Y_4)$
- Think of X as observations and Y as labels

Graphical comparison among HMMs, MEMMs and CRFs




Figure 2. Graphical structures of simple HMMs (left), MEMMs (center), and the chain-structured case of CRFs (right) for sequences. An open circle indicates that the variable is not generated by the model.

Conditional Distribution

If the graph $G = (V, E)$ of Y is a tree, the conditional distribution over the label sequence $Y = y$, given $X = x$, by fundamental theorem of random fields is:

$$p_{\theta}(y | x) \propto \exp \left(\sum_{e \in E, k} \lambda_k f_k(e, y|_e, x) + \sum_{v \in V, k} \mu_k g_k(v, y|_v, x) \right)$$

x is a data sequence

y is a label sequence

v is a vertex from vertex set V = set of label random variables

e is an edge from edge set E over V

f_k and g_k are given and fixed. g_k is a Boolean vertex feature; f_k is a Boolean edge feature

k is the number of features

$\theta = (\lambda_1, \lambda_2, L, \lambda_n; \mu_1, \mu_2, L, \mu_n); \lambda_k$ and μ_k are parameters to be estimated

$y|_e$ is the set of components of y defined by edge e

$y|_v$ is the set of components of y defined by vertex v

Conditional Distribution (cont'd)

- CRFs use the observation-dependent normalization $Z(x)$ for the conditional distributions:

$$p_{\theta}(y|x) = \frac{1}{Z(x)} \exp \left(\sum_{e \in E, k} \lambda_k f_k(e, y|_e, x) + \sum_{v \in V, k} \mu_k g_k(v, y|_v, x) \right)$$

$Z(x)$ is a normalization over the data sequence x

Parameter Estimation for CRFs

- The paper provided iterative scaling algorithms
- It turns out to be very inefficient
- Prof. Dietterich's group applied **Gradient Descendent Algorithm**, which is quite efficient

Training of CRFs (From Prof. Dietterich)

- First, we take the log of the equation

$$\log p_\theta(y | x) = \sum_{e \in E, k} \lambda_k f_k(e, y|_e, x) + \sum_{v \in V, k} \mu_k g_k(v, y|_v, x) - \log Z(x)$$

- Then, take the derivative of the above equation

$$\frac{\partial \log p_\theta(y | x)}{\partial \theta} = \frac{\partial}{\partial \theta} \left(\sum_{e \in E, k} \lambda_k f_k(e, y|_e, x) + \sum_{v \in V, k} \mu_k g_k(v, y|_v, x) - \log Z(x) \right)$$


- For training, the first 2 items are easy to get.
- For example, for each λ_k, f_k is a sequence of Boolean numbers, such as 00101110100111.

$\lambda_k f_k(e, y|_e, x)$ is just the total number of 1's in the sequence.

- The hardest thing is how to calculate $Z(x)$

Training of CRFs (From Prof. Dietterich) (cont'd)

- Maximal cliques



$$c_1 : \exp(\varphi(y_1, x) + \varphi(y_2, x) + \psi(y_1, y_2, x)) = c_1(y_1, y_2, x)$$

$$c_2 : \exp(\varphi(y_3, x) + \psi(y_2, y_3, x)) = c_2(y_2, y_3, x)$$

$$c_3 : \exp(\varphi(y_4, x) + \psi(y_3, y_4, x)) = c_3(y_3, y_4, x)$$


$$\begin{aligned} Z(x) &= \sum_{y_1, y_2, y_3, y_4} c_1(y_1, y_2, x) c_2(y_2, y_3, x) c_3(y_3, y_4, x) \\ &= \sum_{y_1} \sum_{y_2} c_1(y_1, y_2, x) \sum_{y_3} c_2(y_2, y_3, x) \sum_{y_4} c_3(y_3, y_4, x) \end{aligned}$$

Modeling the label bias problem

- In a simple HMM, each state generates its designated symbol with probability $29/32$ and the other symbols with probability $1/32$
- Train MEMM and CRF with the same topologies
- A run consists of 2,000 training examples and 500 test examples, trained to convergence using Iterative Scaling algorithm
- CRF error is 4.6%, and MEMM error is 42%
- MEMM fails to discriminate between the two branches
- CRF solves label bias problem


MEMM vs. HMM

- The HMM outperforms the MEMM




MEMM vs. CRF

- CRF usually outperforms the MEMM



CRF vs. HMM

Each open square represents a data set with $\alpha < 1/2$, and a solid circle indicates a data set with $\alpha \geq 1/2$; When the data is mostly second order ($\alpha \geq 1/2$), the discriminatively trained CRF usually outperforms the HMM



POS tagging Experiments

UPenn tagging task: 45 tags (syntactic), 1M words training

DT NN NN
The asbestos fiber ; crocidolite ; VBZ RB JJ
unusually resilient

IN PRP VBZ DT NNS IN RB JJ NNS
once it enters the lungs ; with even brief exposures

TO PRP VBG NNS WDT VBP RP NNS JJ
to it causing symptoms that show up decades later ;

NNS VBD
researchers said

POS tagging Experiments (cont'd)

- Compared HMMs, MEMMs, and CRFs on Penn treebank POS tagging
- Each word in a given input sentence must be labeled with one of 45 syntactic tags
- Add a small set of orthographic features: whether a spelling begins with a number or upper case letter, whether it contains a hyphen, and if it contains one of the following suffixes: -ing, -ogy, -ed, -s, -ly, -ion, -tion, -ity, -ies
- oov = out-of-vocabulary (not observed in the training set)

<i>model</i>	<i>error</i>	<i>oov error</i>
HMM	5.69%	45.99%
MEMM	6.37%	54.61%
CRF	5.55%	48.05%
MEMM ⁺	4.81%	26.99%
CRF ⁺	4.27%	23.76%

⁺Using spelling features

Summary

- Discriminative models are prone to the label bias problem
- CRFs provide the benefits of discriminative models
- CRFs solve the label bias problem well, and demonstrate good performance

Thanks for your attention!

Special thanks to
Prof. Dietterich & Tadepalli!