

HW4_Solution

February 3, 2024

0.1 Question 1

0.1.1 Dataset

homework4_file1.csv

0.1.2 Data Description

The dataset contains records of merchant transactions, each with a unique merchant identifier, time of transaction, and amount in cents.

0.1.3 Objective

Analyze merchant transaction data to understand business growth and health. Preprocess the dataset for future merchant transactions and generate specific features for each merchant.

0.1.4 Task

Generate the following features for each unique merchant: - **trans__amount__avg**: Average transaction amount for each merchant. - **trans__amount__volume**: Total transaction amount for each merchant. - **trans__frequency**: Total count of transactions for each merchant. - **trans__recency**: Recency of the last transaction (in days from 1/1/2035). - **avg__time__btwn__trans**: Average time between transactions (in hours). - **avg__trans__growth__rate**: Average growth rate in transaction amounts.

0.1.5 Data Dimension

The dataset is N by 3, where N is the number of records.

0.1.6 Final Deliverables

- Shape of the new dataset.
- The top five rows of the new dataset using `new_dataset.head()`.
- Descriptive statistics of the new dataset.

```
[2]: import pandas as pd
import numpy as np
```

```
[3]: data1 = pd.read_csv('homework4_file1.csv')
```

```
[5]: data1.shape
```

```
[5]: (100000, 3)
```

```
[7]: data1.head()
```

```
[7]:      merchant      time  amount_usd_in_cents
0  d087d4c321  2034-12-11 22:16:41             5059
1  fe1cb2e840  2034-08-13 21:11:59            12743
2  878047f4b9  2033-06-05 21:15:00             7601
3  3932608d23  2034-04-28 19:55:01             5790
4  84a09b4188  2034-07-26 04:37:05             6153
```

```
[8]: #Count unique merchant
data1['merchant'].nunique()
```

```
[8]: 7902
```

```
[9]: #transform to datetime format
data1['time'] = data1['time'].apply(pd.to_datetime)
```

```
[10]: #sort time, group by merchant, and aggregate into a list
data1 = data1.sort_values(by='time')
data1_agg = data1.groupby('merchant').agg(list).reset_index()
```

```
[11]: data1_agg.head()
```

```
[11]:      merchant      time \
0  00057d4302  [2033-05-30 02:54:34, 2033-05-30 04:20:31]
1  000ed1585f  [2033-05-08 15:51:43, 2033-05-19 13:15:51, 203...
2  000f8c3297  [2033-12-20 17:57:20, 2034-01-26 15:10:54, 203...
3  0020aefbd9  [2034-05-30 21:55:06]
4  0026f256ac  [2033-09-15 01:17:32]

      amount_usd_in_cents
0  [1156, 1279]
1  [32004, 35784, 21932, 22481]
2  [7374, 14489, 15047, 4623, 3826, 3643, 4155, 7...
3  [3589]
4  [34880]
```

```
[13]: data1_agg.shape
```

```
[13]: (7902, 3)
```

```
[14]: #function that calculates the average time between transactions for each
      ↪merchant
def avg_time_btwn_trans(y):
    if len(y) == 1:
```

```

        return 0
    else:
        td = np.mean([z1 - z0 for z0, z1 in zip(y,y[1:])])
        td_hrs = (td.days)*24 + (td.seconds)/(60*60)
        avg_time_btwn_trans = np.round(td_hrs, 2)
        return avg_time_btwn_trans

```

```

[15]: #function that calculates the average transaction growth rate for each merchant
def avg_trans_growth_rate(y):
    if len(y) == 1:
        return 0
    else:
        grw = np.mean([(z1 - z0)/z0 for z0, z1 in zip(y,y[1:])])
        avg_trans_growth_rate = np.round(grw, 2)
        return avg_trans_growth_rate

```

```

[18]: data1['time'].max()

```

```

[18]: Timestamp('2034-12-31 07:55:58')

```

```

[19]: def get_aggregated(x):
        x['trans_amount_avg'] = x['amount_usd_in_cents'].apply(lambda y: np.
↪round(np.mean(y), 2))
        x['trans_amount_volume'] = x['amount_usd_in_cents'].apply(lambda y: np.
↪round(np.sum(y),2))
        x['trans_frequency'] = x['amount_usd_in_cents'].apply(lambda y: len(y))
        x['trans_recency'] = x['time'].apply(lambda t: (np.datetime64('2035-01-01')
↪- np.max(t)).days + 1)
        x['avg_time_btwn_trans'] = x['time'].apply(lambda t: avg_time_btwn_trans(t))
        x['avg_trans_growth_rate'] = x['amount_usd_in_cents'].apply(lambda y:
↪avg_trans_growth_rate(y))
        return x

```

```

[20]: new_dataset = get_aggregated(data1_agg)

```

```

[21]: new_dataset.head()

```

```

[21]:      merchant                                time \
0  00057d4302      [2033-05-30 02:54:34, 2033-05-30 04:20:31]
1  000ed1585f  [2033-05-08 15:51:43, 2033-05-19 13:15:51, 203...
2  000f8c3297  [2033-12-20 17:57:20, 2034-01-26 15:10:54, 203...
3  0020aefbd9                                [2034-05-30 21:55:06]
4  0026f256ac                                [2033-09-15 01:17:32]

                                amount_usd_in_cents  trans_amount_avg \
0                                [1156, 1279]                1217.50
1          [32004, 35784, 21932, 22481]                28050.25

```

```

2 [7374, 14489, 15047, 4623, 3826, 3643, 4155, 7...      6635.56
3                                     [3589]          3589.00
4                                     [34880]         34880.00

   trans_amount_volume  trans_frequency  trans_recency  avg_time_btwn_trans \
0                2435                2            581            1.43
1             112201                4            175          3424.03
2             106169               16             59           508.47
3              3589                1            216             0.00
4             34880                1            473             0.00

   avg_trans_growth_rate
0                0.11
1             -0.08
2                0.23
3                0.00
4                0.00

```

```
[22]: new_dataset.describe()
```

```

[22]:      trans_amount_avg  trans_amount_volume  trans_frequency  trans_recency \
count      7.902000e+03      7.902000e+03      7902.000000      7902.000000
mean       3.073318e+04      1.963547e+05      12.655024      170.320299
std        1.417803e+05      6.000438e+05      46.531552      180.309019
min        2.090000e+02      2.090000e+02      1.000000      1.000000
25%        4.846177e+03      1.025200e+04      1.000000      26.000000
50%        9.053630e+03      3.484000e+04      3.000000      98.000000
75%        2.114705e+04      1.388630e+05      8.000000      265.000000
max        1.038551e+07      1.549983e+07      1673.000000      727.000000

      avg_time_btwn_trans  avg_trans_growth_rate
count      7902.000000      7902.000000
mean        749.494185        1.011835
std       1461.800362        9.954018
min          0.000000       -1.000000
25%          0.000000        0.000000
50%        170.445000        0.030000
75%        841.512500        0.560000
max       15327.180000       606.650000

```

0.2 Question 2

0.2.1 Datasets Provided

- sales_data.csv
- product_info.csv

0.2.2 sales_data.csv

Contains transaction records with columns: - TransactionID - ProductID - Date - Quantity - Price

0.2.3 product_info.csv

Contains product details with columns: - ProductID - ProductName - Category

0.2.4 Tasks

Your task involves multiple steps of data manipulation using Pandas and NumPy to extract insights from these datasets.

0.2.5 1. Data Loading and Merging

- Load both datasets using Pandas.
- Merge them into a single DataFrame on ProductID.

0.2.6 2. Data Cleaning

- Check for and handle any missing values in the merged dataset.
- Convert the Date column to a DateTime object.

0.2.7 3. Data Analysis using Slicing and Indexing

- Create a new column TotalSale, calculated as Quantity * Price.
- Using slicing, create a subset DataFrame containing only transactions from the last quarter of the year (October, November, December).
- Using Boolean indexing, find all transactions for a specific Category (e.g., Electronics).
- Extract all transactions where the TotalSale is above the 75th percentile of the TotalSale column using NumPy functions.

0.2.8 4. Advanced Indexing

- Using loc and iloc, perform the following:
 - Select all rows for ProductID 101 and columns ProductName and TotalSale.
 - Select every 10th row from the merged dataset and only the columns Date and Category.

0.2.9 5. Grouping and Aggregation

- Group the data by Category and calculate the total and average TotalSale for each category.

0.2.10 6. Time-Series Analysis

- Resample the data on a monthly basis and calculate the total Quantity sold per month.

0.2.11 Final Deliverables

- Provide the code for each step.
- Include comments explaining your approach.

- Display the first 5 rows of the DataFrame after each major step.

```
[23]: # Load datasets
sales_data = pd.read_csv('sales_data.csv')
product_info = pd.read_csv('product_info.csv')
```

```
[24]: # Merge datasets on ProductID
merged_data = pd.merge(sales_data, product_info, on='ProductID')
```

```
[25]: sales_data.head()
```

```
[25]:
```

	TransactionID	ProductID	Date	Quantity	Price
0	1	136	2023-03-13	8	245.288680
1	2	121	2023-06-09	2	355.603776
2	3	179	2023-04-18	7	25.393345
3	4	142	2023-09-03	10	260.758110
4	5	101	2023-06-21	1	212.490775

```
[27]: product_info.head()
```

```
[27]:
```

	ProductID	ProductName	Category
0	100	not	Clothing
1	101	ready	Clothing
2	102	fill	Books
3	103	avoid	Clothing
4	104	beyond	Toys

```
[28]: merged_data.head()
```

```
[28]:
```

	TransactionID	ProductID	Date	Quantity	Price	ProductName	\
0	1	136	2023-03-13	8	245.288680		pull
1	92	136	2023-07-02	6	21.266893		pull
2	260	136	2023-04-15	2	356.242853		pull
3	411	136	2023-08-21	2	91.071146		pull
4	479	136	2023-03-02	10	331.557053		pull

```
Category
```

0	Toys
1	Toys
2	Toys
3	Toys
4	Toys

```
[35]: print("Sales Data Shape:", sales_data.shape)
print("Product Info Shape:", product_info.shape)
print("Merged Data Shape:", merged_data.shape)
```

Sales Data Shape: (10000, 5)

Product Info Shape: (100, 3)
Merged Data Shape: (10000, 7)

```
[36]: # Check for missing values
print(merged_data.isnull().sum())
```

```
TransactionID    0
ProductID        0
Date             0
Quantity         0
Price            0
ProductName       0
Category         0
dtype: int64
```

```
[37]: merged_data['Date'] = pd.to_datetime(merged_data['Date'])
```

```
[47]: # Create 'TotalSale' column
merged_data['TotalSale'] = merged_data['Quantity'] * merged_data['Price']

# Subset for last quarter of the year
last_quarter_data = merged_data[merged_data['Date'].dt.month.isin([10, 11, 12])]

# Boolean indexing for a specific category, e.g., 'Electronics'
electronics_data = merged_data[merged_data['Category'] == 'Electronics']

# Transactions above the 75th percentile of 'TotalSale'
percentile_75 = np.percentile(merged_data['TotalSale'], 75)
high_value_sales_75 = merged_data[merged_data['TotalSale'] > percentile_75].
    ↪reset_index(drop = True)
```

```
[48]: high_value_sales_75.head()
```

```
[48]:   TransactionID  ProductID      Date  Quantity      Price ProductName \
0             479         136 2023-03-02         10  331.557053      pull
1             692         136 2023-09-28          8  494.070419      pull
2             879         136 2023-10-31         10  499.344566      pull
3             996         136 2024-01-14         10  444.985596      pull
4            1309         136 2023-03-31         10  423.493947      pull
```

```
   Category  TotalSale
0     Toys  3315.570534
1     Toys  3952.563354
2     Toys  4993.445658
3     Toys  4449.855959
4     Toys  4234.939470
```

```
[43]: high_value_sales_75.shape
```

```
[43]: (2500, 8)
```

```
[49]: # Selecting specific rows and columns using loc and iloc
productID_101_data = merged_data.loc[merged_data['ProductID'] == 101,
↳ ['ProductName', 'TotalSale']]
every_10th_row = merged_data.iloc[:, :10, merged_data.columns.
↳ get_indexer(['Date', 'Category'])]
```

```
[51]: productID_101_data.shape
```

```
[51]: (98, 2)
```

```
[52]: productID_101_data.head()
```

```
[52]:
```

	ProductName	TotalSale
394	ready	212.490775
395	ready	1331.007870
396	ready	3311.017493
397	ready	1565.745895
398	ready	74.588211

```
[53]: every_10th_row.shape
```

```
[53]: (1000, 2)
```

```
[54]: every_10th_row.head()
```

```
[54]:
```

	Date	Category
0	2023-03-13	Toys
10	2023-10-31	Toys
20	2023-04-21	Toys
30	2023-05-29	Toys
40	2023-12-18	Toys

```
[55]: # Group by 'Category' and calculate total and average 'TotalSale'
grouped_data = merged_data.groupby('Category')['TotalSale'].agg(['sum',
↳ 'mean']).reset_index()

print(grouped_data.shape)
grouped_data.head()
```

```
(5, 3)
```

```
[55]:
```

	Category	sum	mean
0	Books	2.756942e+06	1405.169284

1	Clothing	2.547137e+06	1339.893113
2	Electronics	2.151251e+06	1468.430950
3	Home Appliances	3.339347e+06	1414.378361
4	Toys	3.320096e+06	1436.649185

```
[59]: # Resample data on a monthly basis and calculate total 'Quantity', returning a
      ↪ DataFrame
      monthly_sales = merged_data.resample('M', on='Date').agg({'Quantity': 'sum'}).
      ↪ reset_index()
```

```
[60]: print(monthly_sales.shape)
      monthly_sales.head()
```

(13, 2)

```
[60]:      Date  Quantity
0 2023-01-31      902
1 2023-02-28     4175
2 2023-03-31     4874
3 2023-04-30     4375
4 2023-05-31     4851
```

0.3 Question 3

Zillow's marketplace offers a data-driven home valuation platform utilized by a diverse range of users including home buyers, sellers, renters, homeowners, real estate agents, mortgage providers, property managers, and landlords. The machine learning and data science team at Zillow employs various tools for predicting home valuations, such as Zestimate (Zillow Estimate), Zestimate Forecast, Zillow Home Value Index, Rent Zestimate, Zillow Rent Index, and the Pricing Tool.

0.3.1 Assignment Overview:

You are provided with a dataset named `zillow_feature_sample.csv`, containing various features relevant to Zillow's marketplace. Accompanying the dataset is a data dictionary titled `zillow_data_dictionary.xlsx`, which details the description of each column.

0.3.2 Tasks:

1. *Develop a Missing Data Strategy:* - Assess the `zillow_feature_sample.csv` dataset and devise a comprehensive strategy to handle missing data.
2. *Quantitative Analysis of Missing Data:* - Calculate and report the percentage of missing data in each feature of the dataset. - Analyze and infer the potential mechanism of missing data (e.g., Missing Completely at Random, Missing at Random, Missing Not at Random).
3. *Imputation Strategy:* - Propose and justify an imputation strategy for the missing values in the dataset. Your rationale should be data-driven and well-explained.
4. *Open-Ended Exploration:* - This question is open-ended, allowing you to explore other relevant aspects of the dataset. Conduct additional analyses or apply data processing techniques as

appropriate.

0.3.3 Submission Guidelines:

- Document your analysis and findings in a clear and structured format.
- Ensure that your submission is thorough and well-reasoned.

```
[62]: #!pip install ydata-profiling
```

```
[64]: from ydata_profiling import ProfileReport
```

```
[65]: zillow_data = pd.read_csv("zillow_feature_sample.csv")
```

```
[67]: zillow_data.shape
```

```
[67]: (10000, 58)
```

```
[68]: zillow_data.head()
```

```
[68]:   parcelid  airconditioningtypeid  architecturalstyletypeid  basementsqft  \
0   12833975                    NaN                    NaN            NaN
1   11070096                    1.0                    NaN            NaN
2   12752672                    1.0                    NaN            NaN
3   11338563                    NaN                    NaN            NaN
4   17098704                    NaN                    NaN            NaN

      bathroomcnt  bedroomcnt  buildingclasstypeid  buildingqualitytypeid  \
0              3.0          4.0                    NaN              6.0
1              4.0          4.0                    NaN              7.0
2              2.0          3.0                    NaN              6.0
3              3.0          4.0                    NaN              7.0
4              0.0          3.0                    NaN              NaN

      calculatedbathnbr  decktypeid  ...  numberofstories  fireplaceflag  \
0              3.0          NaN  ...              NaN            NaN
1              4.0          NaN  ...              NaN            NaN
2              2.0          NaN  ...              NaN            NaN
3              3.0          NaN  ...              NaN            NaN
4              NaN          NaN  ...              1.0            NaN

      structuretaxvaluedollarcnt  taxvaluedollarcnt  assessmentyear  \
0              155403.0          304592.0          2016.0
1              493070.0          821783.0          2016.0
2              126695.0          247962.0          2016.0
3              130500.0          308900.0          2016.0
4              142271.0          223101.0          2016.0

      landtaxvaluedollarcnt  taxamount  taxdelinquencyflag  taxdelinquencyyear  \
```

0	149189.0	3708.29	NaN	NaN
1	328713.0	10087.59	NaN	NaN
2	121267.0	3377.86	NaN	NaN
3	178400.0	3578.92	NaN	NaN
4	80830.0	2564.86	NaN	NaN

	censustractandblock
0	6.037409e+13
1	6.037108e+13
2	6.037504e+13
3	6.037920e+13
4	6.111000e+13

[5 rows x 58 columns]

```
[70]: zillow_profile = ProfileReport(zillow_data, title="Zillow Profiling Report")
zillow_profile.to_file("zillow_profile_report.html")
```

```
/Users/wale/anaconda3/lib/python3.10/site-
packages/ydata_profiling/profile_report.py:354: UserWarning: Try running
command: 'pip install --upgrade Pillow' to avoid ValueError
  warnings.warn(
Summarize dataset: 92%|
|
60/65 [00:00<00:00, 49.32it/s, Calculate auto
correlation]/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
  warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
  warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
  warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
  warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
  warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
```

```

constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.
warnings.warn(stats.ConstantInputWarning(warn_msg))
/Users/wale/anaconda3/lib/python3.10/site-
packages/scipy/stats/_stats_py.py:4921: ConstantInputWarning: An input array is
constant; the correlation coefficient is not defined.

```

```
constant; the correlation coefficient is not defined.  
  warnings.warn(stats.ConstantInputWarning(warn_msg))  
Summarize dataset: 100%|
```

```
  | 1364/1364 [01:14<00:00, 18.28it/s, Completed]  
Generate report structure: 100%|
```

```
  | 1/1 [00:06<00:00, 6.21s/it]  
Render HTML: 100%|
```

```
  | 1/1 [00:10<00:00, 10.95s/it]  
Export report to file: 100%|
```

```
  | 1/1 [00:00<00:00, 11.84it/s]
```

[]: