

I probably spent too much time on this assignment this week. You may not see that in the results. I actually enjoyed much of it and learned a lot. In many cases, I just printed the dataframe rather than using head. This provided the first five rows and last five rows. That seemed useful to me. Some of the data is very wide. So, I will upload a text file along with this. It is easier to see the wide data on a text file.

Let me know if you want to see the source file. This is the executed file and has everything in it. But I am happy to share the source file. This assignment is in GitHub at <https://github.com/OwlSaver/GWU>.

## Execution

```
#####
# Problem 1
#####
```

Problem:

Dataset: homework4 file1.csv

Data Description: The dataset contains records of merchant transactions, each with a unique merchant identifier, time of transaction, and amount in cents.

Objective: Analyze merchant transaction data to understand business growth and health. Preprocess the dataset for future merchant transactions and generate specific features for each merchant.

Task: Generate the following features for each unique merchant:

- trans amount avg: Average transaction amount for each merchant.
- trans amount volume: Total transaction amount for each merchant.
- trans frequency: Total count of transactions for each merchant.
- trans recency: Recency of the last transaction (in days from 1/1/2035).
- avg time btwn trans: Average time between transactions (in hours).
- avg trans growth rate: Average growth rate in transaction amounts.

Data Dimension: The dataset is N by 3, where N is the number of records.

Final Deliverables:

- Shape of the new dataset.
- The top five rows of the new dataset using new dataset.head().
- Descriptive statistics of the new dataset.

Code:

```
import pandas as pd
import numpy as np
import datetime as dt
pd.options.display.float_format = '{:,.2f}'.format
pd.set_option('display.max_columns', None)
pd.set_option('display.width', 2000)
HW4F1 = pd.read_csv('./gwu/SEAS 6414/homework4_file1.csv')
# Make the time column a Pandas time rather than a string
HW4F1['time'] = [pd.Timestamp(ts) for ts in HW4F1.time]
HW4F1.sort_values(by=['merchant', 'time'], inplace=True)
HW4F1New = HW4F1.groupby("merchant").agg(
    min_amount=("amount_usd_in_cents", "min")
    , max_amount=("amount_usd_in_cents", "max")
    , trans_amount_avg=("amount_usd_in_cents", "mean")
    , trans_amount_volume=("amount_usd_in_cents", "sum")
)
```

```

, trans_frequency=("amount_usd_in_cents", "count")
, most_recent_date=("time", "max")
, avg_time_btwn_trans=("time", lambda group: group.sort_values().diff().mean().seconds/(60*60))
, avg_trans_growth_rate=("amount_usd_in_cents", lambda group: group.sort_values().pct_change().mean())
)
# I tried to do this as a Lambda in the agg, but it would not recognize the dt library
# So, in the agg, I find the max and here I calculate the delta
HW4F1New['trans_recency'] = (HW4F1New['most_recent_date'] - dt.datetime(2035, 1, 1)).dt.days
# getting rid of the no longer needed maximum value
HW4F1Final = HW4F1New.drop(columns=['most_recent_date'])

```

```

print(f"The shape of the original data frame is: {HW4F1.shape}")
print(f"The shape of the new data frame is: {HW4F1Final.shape}")
print("")
print("The top five rows are:")
print(HW4F1Final.head(5))
print("")
print("Descriptive statistics:")
print(HW4F1Final.describe())

```

Execution:

```

The shape of the original data frame is: (100000, 3)
The shape of the new data frame is: (7902, 8)

```

The top five rows are:

	min_amount	max_amount	trans_amount_avg	trans_amount_volume	trans_frequency	avg_time_btwn_trans
avg_trans_growth_rate						
merchant						
00057d4302	1156	1279	1,217.50	2435	2	1.43
0.11	-581					
000ed1585f	21932	35784	28,050.25	112201	4	16.03
0.19	-175					
000f8c3297	3455	15047	6,635.56	106169	16	4.47
0.12	-59					
0020aefbd9	3589	3589	3,589.00	3589	1	NaN
NaN	-216					
0026f256ac	34880	34880	34,880.00	34880	1	NaN
NaN	-473					

Descriptive statistics:

	min_amount	max_amount	trans_amount_avg	trans_amount_volume	trans_frequency	avg_time_btwn_trans
avg_trans_growth_rate						
count	7,902.00	7,902.00	7,902.00	7,902.00	7,902.00	5,253.00
5,253.00	7,902.00					
mean	20,390.86	55,609.45	30,733.18	196,354.72	12.66	11.45
1.90	-170.32					
std	135,797.59	187,450.07	141,780.27	600,043.78	46.53	7.43
21.23	180.31					
min	201.00	209.00	209.00	209.00	1.00	0.00
0.00	-727.00					
25%	2,061.25	6,585.75	4,846.18	10,252.00	1.00	4.69
0.13	-265.00					
50%	4,226.00	15,442.00	9,053.63	34,840.00	3.00	11.45
0.33	-98.00					
75%	10,510.25	40,685.00	21,147.05	138,863.00	8.00	18.04
0.89	-26.00					
max	10,385,508.00	10,385,508.00	10,385,508.00	15,499,827.00	1,673.00	24.00
1,224.39	-1.00					

```

#####
# Problem 2
#####

```

Problem:

You are provided with two datasets: sales data.csv and product info.csv.  
- sales data.csv contains transaction records with columns: 'TransactionID', 'ProductID', 'Date', 'Quantity', and 'Price'.  
- product info.csv contains product details with columns: 'ProductID', 'ProductName', 'Category'.

Your task involves multiple steps of data manipulation using Pandas and NumPy to extract insights from these datasets.

Tasks:

1. Data Loading and Merging:
  - Load both datasets using Pandas.

- Merge them into a single DataFrame on 'ProductID'.
  - 2. Data Cleaning:
    - Check for and handle any missing values in the merged dataset.
    - Convert the 'Date' column to a DateTime object.
  - 3. Data Analysis using Slicing and Indexing:
    - Create a new column 'TotalSale', calculated as 'Quantity' \* 'Price'.
    - Using slicing, create a subset DataFrame containing only transactions from the last quarter of the year (October, November, December).
    - Using Boolean indexing, find all transactions for a specific 'Category' (e.g., 'Electronics').
    - Extract all transactions where the 'TotalSale' is above the 75th percentile of the 'TotalSale' column using NumPy functions.
  - 4. Advanced Indexing:
    - Using loc and iloc, perform the following:
      - Select all rows for 'ProductID' 101 and columns 'ProductName' and 'TotalSale'.
      - Select every 10th row from the merged dataset and only the columns 'Date' and 'Category'.
  - 5. Grouping and Aggregation:
    - Group the data by 'Category' and calculate the total and average 'TotalSale' for each category.
  - 6. Time-Series Analysis:
    - Resample the data on a monthly basis and calculate the total 'Quantity' sold per month.
- Final Deliverables:
- Provide the code for each step.
  - Include comments explaining your approach.
  - Display the first 5 rows of the DataFrame after each major step.

Code:

```
import numpy as np
import pandas as pd
pd.options.display.float_format = '{:,.2f}'.format
print("Task 1 - Data Loading and Merging")
SalesData = pd.read_csv("./gwu/SEAS 6414/sales_data.csv")
Product = pd.read_csv("./gwu/SEAS 6414/product_info.csv")
# I checked the row counts and there are no product ids in the Sales Data
# that have product keys that are not in Product data. So, an inner join
# will work for this data.
SalesProductData = pd.merge(SalesData, Product, on="ProductID", how="inner")
print("The merged SalesProductData data frame.")
print(SalesProductData)
print("")
print("Task 2 - Data Cleaning")
# Counting the NAs across the dimensions shows that there is no missing data. I also ran
# dropna and saw that the result had the same shape as the input. So, I am confident that
# there is no missing data. Which worries me. Why would you ask us to address missing data
# if there was none.
print(f"The merged data frame has {SalesProductData.isnull().sum().sum()} missing values.")
print("")
print("The SalesProductData types before converting to a datetime:")
print(SalesProductData.dtypes)
SalesProductData['Date'] = [pd.to_datetime(aDate) for aDate in SalesProductData.Date]
print("")
print("The SalesProductData types after converting to a datetime:")
print(SalesProductData.dtypes)
print("")
print("Task 3 - Data Analysis using Slicing and Indexing")
SalesProductData['TotalSale'] = SalesProductData['Quantity'] * SalesProductData['Price']
SalesProductData4Q = SalesProductData.set_index('Date').sort_values(by=['Date'])['2023-10-01' : '2023-12-31']
print("")
print("Sales records for the fourth quarter:")
print(SalesProductData4Q)
mask = SalesProductData['Category'] == 'Electronics'
SalesProductElectronics = SalesProductData[mask]
print("")
print("Sales records for Electronics:")
print(SalesProductElectronics)
# First create and index of all records that have a TotalSale value greater than the 75th percentile
SalesProductOver75Index =
np.where(SalesProductData['TotalSale'] > np.percentile(SalesProductData['TotalSale'], 75))
# Next select those values.
SalesProductOver75 = SalesProductData.loc[SalesProductOver75Index]
print("")
print("Sales records for total price over the 75th percentile:")
print(SalesProductOver75)
```

```

print("")
print("Task 4 - Advanced Indexing")
SalesProductDataPID = SalesProductData.set_index('ProductID')
SalesProductDataI01 = SalesProductDataPID.loc[101,['ProductName','TotalSale']]
print("")
print("Sales records for product 101 with Product Name and Total Sale:")
print(SalesProductDataI01)
SalesProductDataEvery10th = SalesProductData.iloc[::10,[2,6]]
print("")
print("Sales records for every 10th row with Date and Category:")
print(SalesProductDataEvery10th)
print("")
print("Task 5 - Grouping and Aggregation")
SalesProductDataCatGrp = SalesProductData.groupby("Category").agg(
    total_sale=("TotalSale", "sum")
    , average_sale=("TotalSale", "mean")
)
print("")
print("Sales records grouped by category with total and average sales by category:")
print(SalesProductDataCatGrp)
print("")
print("Task 6 - Time-Series Analysis")
# Get down to just the columns needed. I tried to combine this with the indexing but
# none of my incantations would work.
SalesProductDataSmall = SalesProductData.loc[:,['Date','Quantity']]
# To resample, we need the date to be the index
SalesProductDataDate = SalesProductDataSmall.set_index('Date')
# Now we can resample down to Month End and calculate the average
SalesProductDataMonth = SalesProductDataDate.resample('ME').mean()
print(SalesProductDataMonth)

```

Execution:

#### Task 1 - Data Loading and Merging

The merged SalesProductData data frame.

	TransactionID	ProductID	Date	Quantity	Price	ProductName	Category
0	1	136	2023-03-13	8	245.29	pull	Toys
1	2	121	2023-06-09	2	355.60	left	Home Appliances
2	3	179	2023-04-18	7	25.39	according	Books
3	4	142	2023-09-03	10	260.76	hospital	Toys
4	5	101	2023-06-21	1	212.49	ready	Clothing
...	...	...	...	...	...	...	...
9995	9996	136	2023-01-30	7	29.29	pull	Toys
9996	9997	160	2023-05-23	1	96.70	next	Electronics
9997	9998	122	2023-07-14	10	175.15	product	Toys
9998	9999	116	2023-03-25	10	337.27	carry	Home Appliances
9999	10000	186	2024-01-11	5	451.14	increase	Home Appliances

[10000 rows x 7 columns]

#### Task 2 - Data Cleaning

The merged data frame has 0 missing values.

The SalesProductData types before converting to a datetime:

```

TransactionID    int64
ProductID       int64
Date            object
Quantity        int64
Price           float64
ProductName      object
Category        object
dtype: object

```

The SalesProductData types after converting to a datetime:

```

TransactionID    int64
ProductID       int64
Date            datetime64[ns]
Quantity        int64
Price           float64
ProductName      object
Category        object
dtype: object

```

#### Task 3 - Data Analysis using Slicing and Indexing

Sales records for the fourth quarter:

TransactionID	ProductID	Quantity	Price	ProductName	Category	TotalSale
---------------	-----------	----------	-------	-------------	----------	-----------

Date							
2023-10-01	4495	125	9	155.04	instead	Toys	1,395.37
2023-10-01	3289	172	3	334.31	other	Books	1,002.94
2023-10-01	5299	140	6	15.47	condition	Clothing	92.83
2023-10-01	4878	165	9	271.28	others	Home Appliances	2,441.51
2023-10-01	4358	156	6	335.66	create	Toys	2,013.94
...	...	...	...	...	...	...	...
2023-12-31	1495	172	2	72.93	other	Books	145.86
2023-12-31	7296	158	7	227.04	team	Books	1,589.29
2023-12-31	3479	103	2	276.34	avoid	Clothing	552.68
2023-12-31	4827	135	3	158.90	candidate	Books	476.71
2023-12-31	9126	138	2	174.79	collection	Home Appliances	349.58

[2524 rows x 7 columns]

Sales records for Electronics:

	TransactionID	ProductID	Date	Quantity	Price	ProductName	Category	TotalSale
6	7	134	2023-11-06	3	182.18	interview	Electronics	546.54
50	51	164	2023-12-14	9	480.57	energy	Electronics	4,325.13
54	55	166	2023-04-25	5	410.22	group	Electronics	2,051.08
56	57	145	2023-08-21	5	405.02	market	Electronics	2,025.11
66	67	166	2023-04-21	4	447.68	group	Electronics	1,790.74
...	...	...	...	...	...	...	...	...
9977	9978	177	2023-09-29	8	399.54	floor	Electronics	3,196.35
9980	9981	124	2023-11-11	9	236.24	table	Electronics	2,126.18
9981	9982	134	2023-10-04	4	399.44	interview	Electronics	1,597.76
9985	9986	175	2023-09-24	2	156.95	true	Electronics	313.90
9996	9997	160	2023-05-23	1	96.70	next	Electronics	96.70

[1465 rows x 8 columns]

Sales records for total price over the 75th percentile:

	TransactionID	ProductID	Date	Quantity	Price	ProductName	Category	TotalSale
3	4	142	2023-09-03	10	260.76	hospital	Toys	2,607.58
13	14	186	2023-01-31	8	405.01	increase	Home Appliances	3,240.12
14	15	143	2023-11-30	10	293.56	cup	Clothing	2,935.56
17	18	173	2023-08-08	9	458.96	either	Toys	4,130.64
18	19	172	2023-06-03	6	363.44	other	Books	2,180.66
...	...	...	...	...	...	...	...	...
9986	9987	154	2023-10-09	9	497.43	could	Books	4,476.90
9987	9988	169	2023-04-09	8	341.18	everyone	Toys	2,729.46
9988	9989	113	2023-02-17	6	458.28	positive	Clothing	2,749.66
9998	9999	116	2023-03-25	10	337.27	carry	Home Appliances	3,372.71
9999	10000	186	2024-01-11	5	451.14	increase	Home Appliances	2,255.70

[2500 rows x 8 columns]

#### Task 4 - Advanced Indexing

Sales records for product 101 with Product Name and Total Sale:

	ProductName	TotalSale
ProductID		
101	ready	212.49
101	ready	1,331.01
101	ready	3,311.02
101	ready	1,565.75
101	ready	74.59
...	...	...
101	ready	623.85
101	ready	207.41
101	ready	1,348.78
101	ready	2,056.10
101	ready	1,845.12

[98 rows x 2 columns]

Sales records for every 10th row with Date and Category:

	Date	Category
0	2023-03-13	Toys
10	2023-05-16	Home Appliances
20	2023-12-18	Home Appliances
30	2023-12-05	Books
40	2023-04-07	Books
...	...	...
9950	2024-01-12	Home Appliances
9960	2023-03-02	Electronics
9970	2023-10-09	Clothing

```
9980 2023-11-11      Electronics
9990 2023-04-28      Toys
```

[1000 rows x 2 columns]

#### Task 5 - Grouping and Aggregation

Sales records grouped by category with total and average sales by category:

	total_sale	average_sale
--	------------	--------------

Category		
Books	2,756,942.14	1,405.17
Clothing	2,547,136.81	1,339.89
Electronics	2,151,251.34	1,468.43
Home Appliances	3,339,347.31	1,414.38
Toys	3,320,096.27	1,436.65

#### Task 6 - Time-Series Analysis

	Quantity
--	----------

Date	
2023-01-31	5.40
2023-02-28	5.39
2023-03-31	5.60
2023-04-30	5.56
2023-05-31	5.68
2023-06-30	5.48
2023-07-31	5.57
2023-08-31	5.39
2023-09-30	5.36
2023-10-31	5.59
2023-11-30	5.43
2023-12-31	5.59
2024-01-31	5.40

```
#####
# Problem 3
#####
```

Problem:

Zillow's marketplace offers a data-driven home valuation platform utilized by a diverse range of users including home buyers, sellers, renters, homeowners, real estate agents, mortgage providers, property managers, and landlords. The machine learning and data science team at Zillow employs various tools for predicting home valuations, such as Zestimate (Zillow Estimate), Zestimate Forecast, Zillow Home Value Index, Rent Zestimate, Zillow Rent Index, and the Pricing Tool.

#### Assignment Overview:

You are provided with a dataset named `zillow feature sample.csv`, containing various features relevant to Zillow's marketplace. Accompanying the dataset is a data dictionary titled `zillow data dictionary.xlsx`, which details the description of each column.

#### Tasks:

1. Develop a Missing Data Strategy:
  - Assess the `zillow feature sample.csv` dataset and devise a comprehensive strategy to handle missing data.
2. Quantitative Analysis of Missing Data:
  - Calculate and report the percentage of missing data in each feature of the dataset.
  - Analyze and infer the potential mechanism of missing data (e.g., Missing Completely at Random, Missing at Random, Missing Not at Random).
3. Imputation Strategy:
  - Propose and justify an imputation strategy for the missing values in the dataset. Your rationale should be data-driven and well-explained.
4. Open-Ended Exploration:
  - This question is open-ended, allowing you to explore other relevant aspects of the dataset. Conduct additional analyses or apply data processing techniques as appropriate.

#### Submission Guidelines:

- Document your analysis and findings in a clear and structured format.
- Ensure that your submission is thorough and well-reasoned.

#### Code:

```
import numpy as np
import pandas as pd
pd.options.display.float_format = '{:,.2f}'.format
print("")
```

```

print("Task 1 - Develop a Missing Data Strategy")
ZillowFeatureSample = pd.read_csv("../gwu/SEAS 6414/zillow_feature_sample.csv")
print("The data provided:")
print(ZillowFeatureSample)
print("Descriptive statistics for each feature:")
print(ZillowFeatureSample.describe())
print("")
print("This data is used to predict house prices. Since it does not have actual prices, we cannot")
print("use it for training or testing our models. Therefore, we cannot test the impact of any")
print("missing data strategy with just this data at hand. However, we can look at the data and")
print("determine if any missing data approach would be useful. Below is my strategy based on a")
print("review of the data values and data dictionary.")
print("")
print("From the data dictionary:")
print(" - The data dictionary has eight tabs.")
print(" - The first one is for the data file.")
print(" - The remaining seven are code tables for features that are coded.")
print(" - Eight of feature descriptions had the phrase 'if any' in them, or should.")
print(" - Some features probably should include 'if any' in the description")
print(" - For example, 'airconditioningtypeid' is described as 'Type of cooling system")
print("   present in the home (if any)'"")
print(" - For example, 'assessmentyear' is described as 'The year of the property tax assessment'."")
print("   Since a house may never have been assessed, this is similar to 'if any'."")
print(" - In both these cases, any unavailable information could be treated as a No or whatever")
print("   is appropriate.")
print(" - Seventeen of the features have the characters ID at the end of the name.")
print(" - Of these seven have tables on other tabs and ten do not.")
print(" - Assignment of an ID means that a process was followed to code the data.")
print(" - Given this process, I would be reluctant to replace the missing data with a value.")
print(" - Some data is dependant on other data.")
print(" - If 'regionidzip' is available, we could use that to fill in City, State, etc.")
print(" - For each feature, we can look into any dependencies that could help derive the values.")
print(" - We will need to be careful with this. We will have to determine the dependencies, then")
print(" - derive the data, then remove the dependant values so that only one of them remains. This")
print(" - ensures that we are only left with independent variables (features).")
print(" - There appear to be a lot of missing values. We will need to carefully consider these")
print("   features. We may need to drop those that are missing too many values.")
print("")
print("Task 2 - Quantitative Analysis of Missing Data")
missing_value_analysis = pd.DataFrame({'count_missing': ZillowFeatureSample.isna().sum(),
                                      , 'percent_missing': ZillowFeatureSample.isnull().sum() * 100 /
len(ZillowFeatureSample)})
print("")
print("Count and percent missing for each feature, sorted low to high by percent:")
print(missing_value_analysis.sort_values(by=['percent_missing']))
print("")
print("Searching the web, it looks like a lot of people consider between 10 and 20% missing")
print("a cutoff point -> more than 20% missing, do not use the feature. But this is always followed")
print("with - there is no hard cutoff point. Since we have 9.25% missing and then 34.00% missing")
print("my working assumption for now is that this will be the cutoff point. But I will continue")
print("analyzing the data to see if some of the features with 34.00% or greater missing are useful.")
print("")
print("Trying to infer the mechanism of missing data will be tricky for me. There are several")
print("reasons for this:")
print(" - I do not know how any of the data was collected.")
print(" - This is not an area that I have any expertise in.")
print("")
print("With those caveats in mind, here is my estimation for each feature.")
print(" - For the 23 features that have a missing percent under 4%, I deem them as not")
print("   really missing. If a value is needed for them, it can easily be imputed.")
print(" - For the 26 features with a missing percent over 70%, I deem them as too much")
print("   missing. I would be hard pressed to impute these values. There may be special")
print("   cases as the analysis progresses.")
print(" - The remaining nine features need to be addressed.")
print(" - Based on the information provided, I cannot say if they are MCAR, MAR, or MNAR.")
print("   I would need details about how the information was collected and about housing")
print("   data.")
print("")
print("Based on the above, I created the table below for values that could be imputed:")
print("   finishedsquarefeet12   Impute from Calculated square feet")
print("   lotsizesquarefeet       Impute from address")
print("   unitcnt                 Do not impute - I expect number of units to be unique")
print("   propertyzoningdesc      Impute from address")
print("   buildingqualitytypeid   Do not impute - an ID")
print("   heatingorsystemtypeid   Do not impute - an ID")
print("   regionidneighborhood    Impute from address")
print("   garagearcnt             Impute from address")

```





NaN	1,957.00	NaN	NaN	126,695.00	247,962.00	2,016.00
121,267.00	3,377.86	NaN	NaN	NaN 60,375,038,011,001.00		
3	11338563	NaN	NaN	NaN	3.00	4.00
NaN	7.00	3.00	NaN	NaN	NaN	
2,280.00	2,280.00	NaN	NaN	NaN	NaN	
NaN 6,037.00	NaN	3.00	NaN	NaN	NaN	NaN
34,674,776.00	-118,418,589.00	856,438.00	NaN	NaN	NaN	NaN
0700	263.00	LCRA10000-	NaN	60,379,201.02	25,468.00	3,101.00
NaN	97,316.00	0.00	NaN	NaN	NaN	1.00
NaN	2,006.00	NaN	NaN	130,500.00	308,900.00	2,016.00
178,400.00	3,578.92	NaN	NaN	NaN 60,379,201,022,009.00		
4	17098704	NaN	NaN	NaN	0.00	3.00
NaN	NaN	NaN	NaN	NaN	1,200.00	
1,200.00	1,200.00	NaN	NaN	NaN	1,200.00	
NaN 6,111.00	NaN	NaN	2.00	400.00	NaN	NaN
34,364,318.00	-119,055,992.00	6,750.00	NaN	NaN	NaN	NaN
1110	261.00	NaN	NaN	61,110,004.00	26,965.00	2,061.00
NaN	97,113.00	5.00	NaN	NaN	NaN	NaN
NaN	1,987.00	1.00	NaN	142,271.00	223,101.00	2,016.00
80,830.00	2,564.86	NaN	NaN	NaN 61,110,004,003,032.00		
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
...	...	...	...	...	...	...
9995	11176537	1.00	NaN	NaN	3.00	4.00
NaN	8.00	3.00	NaN	NaN	NaN	NaN
3,170.00	3,170.00	NaN	NaN	NaN	NaN	NaN
NaN 6,037.00	NaN	3.00	NaN	NaN	NaN	2.00
34,473,832.00	-118,633,360.00	7,341.00	1.00	NaN	NaN	NaN
0104	261.00	LCA22*	NaN	60,379,201.16	10,734.00	3,101.00
NaN	96,398.00	0.00	NaN	NaN	NaN	1.00
NaN	2,000.00	NaN	NaN	259,395.00	399,915.00	2,016.00
140,520.00	5,692.80	NaN	NaN	NaN 60,379,201,162,006.00		
9996	12544738	NaN	NaN	NaN	NaN	1.00
NaN	6.00	1.00	NaN	NaN	NaN	1.00
882.00	882.00	NaN	NaN	NaN	NaN	NaN
NaN 6,037.00	NaN	1.00	NaN	NaN	NaN	7.00
33,767,900.00	-118,166,000.00	16,410.00	NaN	NaN	NaN	NaN
010E	266.00	LBR2N	NaN	60,375,768.01	46,298.00	3,101.00
416,302.00	96,237.00	0.00	NaN	NaN	NaN	1.00
NaN	NaN	1,957.00	NaN	NaN	73,738.00	98,658.00
2,016.00	24,920.00	1,242.21	NaN	NaN	NaN 60,375,768,013,008.00	
9997	12546936	NaN	NaN	NaN	NaN	2.00
NaN	8.00	2.00	NaN	NaN	NaN	2.00
1,308.00	1,308.00	NaN	NaN	NaN	NaN	NaN
NaN 6,037.00	NaN	2.00	NaN	NaN	NaN	2.00
33,765,200.00	-118,177,000.00	28,576.00	NaN	NaN	NaN	NaN
010E	266.00	LBPD5	NaN	60,375,766.01	46,298.00	3,101.00
416,302.00	96,236.00	0.00	NaN	NaN	NaN	1.00
NaN	NaN	1,958.00	NaN	NaN	130,500.00	520,000.00
2,016.00	389,500.00	6,214.01	NaN	NaN	NaN 60,375,766,011,005.00	
9998	10858309	1.00	NaN	NaN	NaN	2.00
NaN	7.00	2.00	NaN	NaN	NaN	2.00
1,250.00	1,250.00	NaN	NaN	NaN	NaN	NaN
NaN 6,037.00	NaN	2.00	NaN	NaN	NaN	2.00
34,150,700.00	-118,441,000.00	23,518.00	1.00	NaN	NaN	NaN
010C	266.00	LARD1.5	NaN	60,371,412.01	12,447.00	3,101.00
27,080.00	96,424.00	0.00	NaN	NaN	NaN	1.00
NaN	NaN	1,979.00	NaN	NaN	112,656.00	167,805.00
2,016.00	55,149.00	2,109.45	NaN	NaN	NaN 60,371,412,012,001.00	
9999	12617699	NaN	NaN	NaN	NaN	2.00
NaN	6.00	2.00	NaN	NaN	NaN	4.00
1,581.00	1,581.00	NaN	NaN	NaN	NaN	NaN
NaN 6,037.00	NaN	2.00	NaN	NaN	NaN	2.00
33,800,022.00	-118,218,562.00	6,444.00	NaN	NaN	NaN	NaN
0100	261.00	LBR1N	NaN	60,375,727.00	46,298.00	3,101.00
275,989.00	96,244.00	0.00	NaN	NaN	NaN	1.00
NaN	NaN	1,977.00	NaN	NaN	101,845.00	129,147.00
2,016.00	27,302.00	1,689.56	NaN	NaN	NaN 60,375,727,004,003.00	

[10000 rows x 58 columns]  
Descriptive statistics for each feature:

Page 10

275.00	61,110,091.00	396,556.00	3,101.00	764,167.00	399,675.00	14.00
7.00	2.00	6.00	7.00	1,260.00	1,123.00	2,015.00
4.00	5,275,190.00	19,310,938.00	2,016.00	14,217,944.00	224,000.93	
15.00	61,110,091,003,005.00					

This data is used to predict house prices. Since it does not have actual prices, we cannot use it for training or testing our models. Therefore, we cannot test the impact of any missing data strategy with just this data at hand. However, we can look at the data and determine if any missing data approach would be useful. Below is my strategy based on a review of the data values and data dictionary.

From the data dictionary:

- The data dictionary has eight tabs.
  - The first one is for the data file.
  - The remaining seven are code tables for features that are coded.
- Eight of feature descriptions had the phrase 'if any' in them, or should.
  - Some features probably should include 'if any' in the description
  - For example, 'airconditioningtypeid' is described as 'Type of cooling system present in the home (if any)'
  - For example, 'assessmentyear' is described as 'The year of the property tax assessment'. Since a house may never have been assessed, this is similar to 'if any'.
  - In both these cases, any unavailable information could be treated as a No or whatever is appropriate.
- Seventeen of the features have the characters ID at the end of the name.
  - Of these seven have tables on other tabs and ten do not.
  - Assignment of an ID means that a process was followed to code the data.
  - Given this process, I would be reluctant to replace the missing data with a value.
- Some data is dependant on other data.
  - If 'regionidzip' is available, we could use that to fill in City, State, etc.
  - For each feature, we can look into any dependencies that could help derive the values.
  - We will need to be careful with this. We will have to determine the dependencies, then derive the data, then remove the dependant values so that only one of them remains. This ensures that we are only left with independent variables (features).
- There appear to be a lot of missing values. We will need to carefully consider these features. We may need to drop those that are missing too many values.

## Task 2 - Quantitative Analysis of Missing Data

Count and percent missing for each feature, sorted low to high by percent:

	count_missing	percent_missing
parcelid	0	0.00
fips	13	0.13
propertylandusetypeid	13	0.13
rawcensustractandblock	13	0.13
regionidcounty	13	0.13
longitude	13	0.13
roomcnt	13	0.13
bedroomcnt	13	0.13
bathroomcnt	13	0.13
assessmentyear	13	0.13
latitude	13	0.13
propertycountylandusecode	14	0.14
regionidzip	42	0.42
taxamount	66	0.66
taxvaluedollarcnt	119	1.19
structuretaxvaluedollarcnt	144	1.44
calculatedfinishedsquarefeet	149	1.49
yearbuilt	166	1.66
regionidcity	210	2.10
landtaxvaluedollarcnt	210	2.10
censustractandblock	240	2.40
fullbathcnt	388	3.88
calculatedbathnbr	388	3.88
finishedsquarefeet12	859	8.59
lotssquarefeet	925	9.25
unitcnt	3400	34.00
propertyzoningdesc	3411	34.11
buildingqualitytypeid	3530	35.30
heatingorsystemtypeid	3757	37.57
regionidneighborhood	6078	60.78
garagecarcnt	6978	69.78
garagetotalsqft	6978	69.78
airconditioningtypeid	7219	72.19
numberofstories	7655	76.55
poolcnt	8162	81.62
pooltypeid7	8275	82.75
threequarterbathnbr	8929	89.29

fireplacecnt	8953	89.53
finishedfloorlsquarefeet	9305	93.05
finishedsquarefeet50	9305	93.05
finishedsquarefeet15	9388	93.88
yardbuildingsqft17	9746	97.46
taxdelinquencyflag	9816	98.16
taxdelinquencyyear	9816	98.16
hashottuborspa	9827	98.27
pooltypeid2	9890	98.90
poolsizeum	9894	98.94
finishedsquarefeet6	9928	99.28
decktypeid	9932	99.32
pooltypeid10	9937	99.37
buildingclasstypid	9961	99.61
finishedsquarefeet13	9974	99.74
typeconstructiontypeid	9980	99.80
architecturalstyletypeid	9987	99.87
yardbuildingsqft26	9988	99.88
fireplaceflag	9989	99.89
basementsqft	9996	99.96
storytypeid	9996	99.96

Searching the web, it looks like a lot of people consider between 10 and 20% missing a cutoff point -> more than 20% missing, do not use the feature. But this is always followed with - there is no hard cutoff point. Since we have 9.25% missing and then 34.00% missing my working assumption for now is that this will be the cutoff point. But I will continue analyzing the data to see if some of the features with 34.00% or greater missing are useful.

Trying to infer the mechanism of missing data will be tricky for me. There are several reasons for this:

- I do not know how of the data was collected.
- This is not an area that I have any expertise in.

With those caveats in mind, here is my estimation for each feature.

- For the 23 features that have a missing percent under 4%, I deem them as not really missing. If a value is needed for them, it can easily be imputed.
- For the 26 features with a missing percent over 70%, I deem them as too much missing. I would be hard pressed to impute these values. There may be special cases as the analysis progresses.
- The remaining nine features need to be addressed.
- Based on the information provided, I cannot say if they are MCAR, MAR, or MNAR. I would need details about how the information was collected and about housing data.

Based on the above, I created the table below for values that could be imputed:

finishedsquarefeet12	Impute from Calculated square feet
lotsizesquarefeet	Impute from address
unitcnt	Do not impute - I expect number of units to be unique
propertyzoningdesc	Impute from address
buildingqualitytypeid	Do not impute - an ID
heatingorsystemtypeid	Do not impute - an ID
regionidneighborhood	Impute from address
garagecarcnt	Impute from address
garagetotalsqft	Impute from address

### Task 3 - Imputation strategy

Let me start by saying that my gut reaction is that using imputation is a really bad idea. We have data that we are trying to use to predict something and before we do we are predicting values that are missing from the data. If we use existing values to impute the values, we are not adding anything to the data we have. I am actually concerned that people are making decisions based on this. It seems like an incredibly bad idea.

If I had to impute values for this data set, I would use averages in most cases. I would try to find a set of the data from the same general area and similar houses. This is based on the idea that all 3,000 square foot houses built in the same area in the same time period will essentially be the same. So, if we can get enough records, we can do that. This data set may be too small to get enough records. But given that Zillow seems to have data for every house in the US, it should be possible to get more data.

Based on this, I would be willing to impute values for the 23 features that are missing under 4% of the values and the four features identified above.

### Task 4 - Open-Ended Exploration

Does year built correlate with size?

yearbuilt	calculatedfinishedsquarefeet
-----------	------------------------------

yearbuilt	1.00	0.17
calculatedfinishedsquarefeet	0.17	1.00

It appears to have a low correlation.

Does latitude correlate air conditioning?

	latitude	airconditioningtypeid
latitude	1.00	-0.45
airconditioningtypeid	-0.45	1.00

This seems to be saying that there is an inverse relation. That makes sense. The higher the latitude, the less need there is for air conditioning. Note that the values for air conditioning are not really good for this correlation. To really do it right, I would need to convert the values. But as a first cut, it makes sense.

I could probably do similar things for pools at lower latitudes and fire places at higher latitudes. I am not sure it would be worthwhile given the amount of missing data.